FINMANAI · PROGRESS REPORT

The Product

What has been built, what it does for a household, and what is still missing before anyone can use it.

> FinManAI takes a complete Indian household financial profile — income, spending, dependants, assets, debts, insurance, goals, risk appetite, tax position — and returns a written, reasoned plan: what to fix first, how much to invest each month, into which asset classes, for which goal. Seven phases are complete. The engine, the security core, the web application, the compliance record and the review cycle all exist and pass their tests. What remains is almost entirely outside the code.

REPORT DATE	ENGINE VERSION	ASSUMPTIONS	ADVISORY MODE	PHASES COMPLETE
10 September 2026	1.0.0	2026-08	diagnostic_only	0–6 of 6

1 Where the product stands

Every phase from design through operations is built and verified. The product can take a real person from a signup form to a stored, explained, printable financial plan, remind them a year later that it has gone stale, show them what changed when they redo it, and give an adviser a queue of who needs attention this month. All of that runs today.

7	475	14	26	0
PHASES COMPLETE	OFFLINE CHECKS PASSING	SCREENS	API ENDPOINTS	REAL USERS YET

Offline checks re-run for this report on 10 September 2026: 103 security-core, 297 engine invariants, 16 narrative-guard, 9 golden snapshots, 50 living-advice checks. All passing.

PHASE	WHAT IT DELIVERS TO THE USER	STATE TODAY	VERIFIED BY
0 Design	Nothing visible. It decided what the product legally is.	COMPLETE The regulatory position was settled before the data model, which is why record-keeping is in the schema rather than bolted on.	Retrospective document
1 Engine	The advice itself — ratios, sequencing, goal maths, allocation, tax, stress tests.	COMPLETE Ten deterministic stages. Same profile in, byte-identical plan out. Nine bugs found and fixed by reading output, not by types.	322 checks · 9 golden snapshots
2 Identity & data	Signing up, two-factor login, and an eight-step profile you can abandon and come back to.	COMPLETE Profiles are immutable versions with a content hash. Encrypted per-field, per-user.	103 checks · driven over HTTP
3 The plan	The plan page, the dashboard, and the printable version a client keeps.	COMPLETE Advisory-mode gate is real: in diagnostic mode the allocation table is genuinely absent, not hidden with CSS.	HTTP suite · gate consistency asserted
4 Narrative	The plan explained in plain language, with a line saying whether a model or a template wrote it.	COMPLETE, UNPROVEN LIVE Guards are tested against stubs. Never run against a live model, so prose quality is unmeasured.	16 checks
5 Compliance	Invisible to the client; decisive in a dispute. Consent, agreement, hash-chained audit trail.	COMPLETE Console at `/compliance`. Tampering is not prevented — it is made provable.	Chain verification · 3 tamper modes
6 Living advice	Annual re-profiling, plan comparison, adviser review queue, reminders.	COMPLETE A plan past its review date is relabelled as a historical record, never withdrawn.	50 checks · 25 of the HTTP checks

> THE HONEST SUMMARY
>
> The software is finished to a pilot standard. The two things standing between it and a real user are not engineering tasks: SEBI Registered Investment Adviser registration before personalised allocation advice may be shown, and a securities lawyer confirming the structure. Until then the product runs in `diagnostic_only` and charges nothing, because any fee is “consideration” and consideration is what makes the advice regulated.

2 What a household actually experiences

The journey below is built end to end. Each step names the screens and the product decision that shaped it — because in a product like this, almost every design choice is a decision about whether someone finishes or gives up.

Step 1 — Sign up and prove it is you

Email and password, or Google sign-in. Then email verification, then TOTP two-factor enrolment with a QR code. **No financial plan is generated for an account without 2FA.** Login is two stages and the first issues nothing: a correct password returns a five-minute challenge ticket, not a session. A stolen password on its own never yields an authenticated session.

Step 2 — The eight-step profile

Personal details, employment and income, spending, dependants, assets and liabilities, insurance and emergency fund, goals and risk questionnaire, then liquidity, tax and consent. Four product rules govern it:

RULE	WHY IT EXISTS
Autosave every 2.5 seconds	People complete this across two sittings. Losing forty minutes of a financial profile to a dropped connection is how you lose the user permanently.
Only five sections are genuinely required	The engine handles absence correctly. A form demanding forty fields before it does anything gets abandoned before it does anything.
Every intrusive question explains itself inline	Asking a child's gender carries the Sukanya Samriddhi reason. Employer health cover says it counts at half weight because it ends with the job. Abandonment happens at exactly the question that feels nosy and unexplained.
Spending is shown, not asked	The Spending step is inverted: it presents a plausible household spending sketch and asks you to correct what is wrong. Recall is a hard task and households underestimate by 20–30%; recognition is easier and more accurate. Every pre-filled figure is labelled as an estimate until somebody disagrees with it, because suitability rests on client-provided information — and a number the client merely failed to contradict is not that.

Step 3 — The plan

Committing the profile writes an immutable version with a SHA-256 of its contents, then runs the engine and stores the result. The plan page is ordered as advice, not as a data dump: the plain-language explanation first, then where the money goes, then health, risk, *what to do in order*, goals, allocation, the existing portfolio, stress tests, tax, and the disclosures — on the plan itself, not behind a footer link.

Step 4 — A year later

Two clocks run separately, deliberately. The **client clock** is the annual review. The **rules clock** is whether the tax slabs and rates the plan was built on are still the current ones — a Budget makes a correct plan stale without the client doing anything. Conflating them produces the sentence every client mistrusts: *your plan has been updated*, with no reason attached.

Reminders go out 30 days before, on the day, then monthly — monthly rather than weekly because the honest message after the first month is not *urgent*, it is *your plan is no longer current advice*, and that is a thing to say once a month, calmly, until it is fixed. Only clients who consented to reminders are mailed, checked at send time because that consent is revocable.

The review screen shows the five inputs that most move the advice and offers two buttons: *everything above is still true*, or *something has changed*. Confirming is one click — a review flow that makes confirming as much work as changing is one people stop completing, and then the record says nothing. And it refuses to accept "nothing has changed" when the saved draft plainly disagrees.

> THE DECISION THAT MOST PROTECTS THE CLIENT
>
> Past the 90-day grace period a stale plan is **relabelled, not withdrawn**. It still renders in full, under a banner saying it is a record of what was advised at the time. Pulling someone's plan the morning it turns a year old — while they may be halfway through acting on it — punishes the client for our reminder schedule. What we told them in March stays exactly what we told them in March.

3 What the advice looks like

Below is genuine engine output for Rohan, 29 — one of nine test households. Every figure is produced by deterministic code. It is included because the character of the product is not visible in a feature list.

Where his money goes

MONTHLY	AMOUNT
Gross income	₹1,08,500
Less: bonus haircut (50%, not planned on)	–₹6,250
Essential, discretionary and irregular spending	–₹56,000
EMIs	–₹22,500
Investable surplus	**₹23,750**

Note the bonus haircut. Variable income is discounted by half before anything is planned on it, because a plan built on a bonus that does not arrive is a plan that fails in the one month it was needed most.

What to do, in order — the sequencing is the advice

	ACTION	MONTHLY	THE REASONING GIVEN TO HIM
1	Buy ₹2.61 Cr of term life cover	₹1,964	He holds none, has a dependent spouse and ₹28.85 L of liabilities. Larger than any market risk, and the cheapest item here — which is why it is funded *before* debt repayment, not after.
2	Add ₹8.5 L of personal health cover	₹1,032	His only cover is his employer's, and it ends the day the job does — which is exactly when he needs it.
3	Clear the ₹85,000 credit card at 42%	₹21,103	A guaranteed, tax-free 42% return. Nothing available in India reliably beats it. Clears in about five months, after which ₹21,103 a month comes back permanently.
4	Top the emergency fund to ₹3.81 L	₹2,490	Not an investment — it is what stops him selling investments at the worst possible moment.
5	Surrender the ULIP when the lock-in ends	—	Worth roughly ₹5.2 L more over twelve years, with the same life cover throughout. It also tells him to get the actual surrender value in writing, because ours is an estimate and it is the one number that changes the answer.
6	Home loan at 8.6% — prepay only if you want to	—	Mathematically, investing usually wins over 15+ years. Emotionally, being debt-free is worth a lot to some people. The product says this is genuinely his call and assumes neither.

7	Invest ₹1,660 now, rising to ₹25,254 in about six months	₹1,660 → ₹25,254	Only ₹1,660 is free today. When the commitments above finish, redirect the released money into the plan rather than into lifestyle — that habit is worth more than any fund selection.

WHERE THE PRODUCT REFUSES TO FLATTER

All three of Rohan's goals come back **Needs More Saving**, and the product says so directly. His retirement goal needs ₹11.3 Cr and only ₹13,168 a month can be allocated after higher-priority commitments. So it gives him the four real options — fifteen more years, a target reduced to about 20% of what he asked for, a larger surplus, or lowering the priority of a competing goal — and closes with: *"Choose one — the goal will not fund itself by being left on the list."* No underfunded goal is ever allowed to fail silently.

Two more things visible in the output

Every allocation line carries its own reason. Small-cap funds are capped at 15% of equity "because small-caps can fall 50–60% and stay down for years". Gold ETFs are preferred over jewellery because of storage, purity and making-charge losses — and the plan notes that Sovereign Gold Bonds would have been better, but no new tranches are being issued.

Recommendations are category level, never named schemes. "Broad-market index fund (Nifty 500 / total market)", not a fund house. That is defensible without RIA registration and carries no data-vendor dependency, and it is a decision that will not be revisited lightly.

4 The decisions that define the product

These are settled. Each was contested, each has a reason, and re-opening one changes what the product is rather than how it is built.

QUESTION	ANSWER	WHY IT IS NOT UP FOR DEBATE
How specific is a recommendation?	**Category level** — "large-cap index fund", never a scheme name	Defensible without RIA registration; no dependency on a market-data vendor.
How is property handled?	A **goal with an affordability model**, not a slice of the portfolio	Illiquid, indivisible, leveraged, 7–10% transaction costs. REITs carry the sector exposure instead.
Crypto?	Opt-in, hard-capped at 5%, modelled at **0% expected return**	Nobody can honestly forecast one. Pretending otherwise is the dishonest part, not the volatility.
Retirement?	A **spending stream to age 85**, tranched by when each rupee is drawn	Modelled as a single date, it put a 57-year-old's entire corpus into short-duration debt — money he would not touch until 75.
Contributions?	**Step-up SIPs**, capped at 10%/yr, 75% of stated income growth	Flat thirty-year SIPs marked almost every goal impossible. The remaining 25% is absorbed by lifestyle, which is what actually happens.
Can a profile be edited?	**No — immutable versions**	A plan generated in March must be explainable in September. If the inputs are overwritten it cannot be.
Can a plan be corrected?	**No — append-only**	Wrong output means a new plan. An audit needs to see what was actually said, not what we wish had been said.
What does the MVP charge?	**Zero**	Any fee is "consideration", and consideration is what makes the advice regulated. Enforced as a hard rejection, not a setting.

The advisory-mode gate, in product terms

`ADVISORY_MODE` has two settings and defaults to the safe one. In `diagnostic_only` the engine still computes the complete plan — but the page withholds the allocation table and the rebalancing actions entirely, and explains why in a banner. Turning advisory output on has to be a deliberate act; it is never the consequence of a missing environment variable.

WHAT THE CLIENT SEES	DIAGNOSTIC_ONLY	ADVISORY
Cashflow, ten financial-health ratios, risk profile	SHOWN	SHOWN
What to do, in order — the action plan	SHOWN	SHOWN
Goal gaps, feasibility and the negotiation	SHOWN	SHOWN
Stress tests — job loss, market crash, medical event	SHOWN	SHOWN
Old vs new tax regime comparison	SHOWN	SHOWN
The allocation table — what to buy, how much	WITHHELD	SHOWN
Existing-portfolio rebalancing actions	WITHHELD	SHOWN
Any fee at all	REJECTED	Capped: ₹1,51,000/family/yr or 2.5% of AUA

Stress tests appear in both modes on purpose. Telling someone they could survive 2.3 months without income is diagnostics, not a recommendation to buy anything.

5 The surfaces beyond the client

The adviser console

In advisory mode the engine's output is not automatically the adviser's advice — an RIA is personally accountable for what goes out under their name, so a person signs off before a plan is released. Approvals are append-only, so a withheld-then-approved sequence stays visible; a mutable column would erase the fact that anyone hesitated. Withholding requires a written reason, because *"why did you not release this?"* is exactly what an audit asks four years later.

The queue sorts clients by urgency — overdue reviews first, then plans awaiting release, then profiles built on superseded assumptions — and warns as the client-count deposit threshold *approaches*, since the IAASB lien has to be arranged before the 151st client rather than discovered after. Adviser notes are encrypted, retained with the plan for the full five years, and never shown to the client: a note written knowing the client is reading over your shoulder is a worse record than one written honestly.

The compliance console

Chain integrity first, because it is the headline an auditor looks for. Then event and plan counts, any plan whose hash no longer matches its profile, the AI-usage breakdown, the advisory-mode split, and the event log with subject and actor shown separately. It renders no plan or profile contents at all — the export deliberately excludes them, and a console that reintroduced what the API refused to ship would defeat the point. An auditor needs to see that advice was given, when, and on what basis; not to receive every client's decrypted financial position in a JSON file.

The plan comparison

Two plans side by side are the plan responding to a life. Three rules shape the view:

1. **Attribute the change first.** "Your SIP went up ₹3,000" is a fact; "…because your income rose and the surplus followed it" is the advice. The cause is named before any number.
2. **Suppress noise.** A ₹40 move in a ₹90,000 figure is arithmetic, not news. The materiality floor is ₹100 for money, half a percentage point for allocation.
3. **Direction is not sentiment.** A larger required SIP is not "worse" — it may be a bigger goal, deliberately chosen. The view reports direction and leaves judgement to the prose.

Where either plan was diagnostic-only, the allocation table is flagged as *withheld* rather than diffed. Reporting a regulatory gate as a change in advice would be actively misleading.

6 What is not built

Listed by what a user would notice, not by what is technically outstanding. Nothing here blocks a pilot; several items would materially improve it.

ITEM	STATUS	WHAT IT MEANS FOR THE PRODUCT
RIA registration	BLOCKING	Required before `ADVISORY_MODE=advisory`. Until it lands, no client sees an allocation table and no fee may be charged. This is the single item that most limits the product, and it is not an engineering task.
Legal review	BLOCKING	A securities lawyer confirming the pilot structure. Everything regulatory in the repository is one reading of the regulations, not legal advice.
Live market-data feeds	STAGED	The discipline is built — stage, review, apply, writing a new dated assumptions version and never editing one in place. The network adapters (AMFI NAV, gold spot, MoSPI CPI) are not; they need credentials and endpoints. Today every rate is updated by hand, which is safe but does not scale.
The narrative against a live model	UNTESTED	Every guard is tested against stubs and works. Nobody has read a real model's prose about a real plan. The guards are proven; the writing is not.
Monte Carlo for retirement	DEFERRED	Three deterministic stress tests are honest but coarse. A sequence-of-returns simulation would strengthen the retirement model — the deferral is deliberate, since a probability distribution invites false precision.
Behavioural nudges	DEFERRED	The largest determinant of outcomes is whether someone keeps investing through a fall. Nothing in the product addresses that yet. This is arguably the biggest missing piece of *product*, as distinct from missing engineering.
Rent vs buy modelling	DEFERRED	The house-purchase goal has an affordability model but does not answer the question most Indian households are actually asking.
Hindi and regional languages	DEFERRED	The narrative layer can do it. The numeric guards are language-agnostic; the claim patterns are English-only and would need rewriting first. Decide before the pilot, not during it.
Scheme-level recommendations	CONDITIONAL	Only if RIA registration lands *and* a documented, rule-based selection methodology exists. Category level remains the safer default and may remain the permanent one.

> ONE DOCUMENTATION GAP WORTH NAMING
>
> The phase documents record 36 living-advice checks; the suite now runs 50, having grown with the Google sign-in work that postdates the phase write-ups. The tests are ahead of the documentation, which is the right direction for the drift to run, but the phase documents should be reconciled before they are shown to anyone external.

7 Where the product risk actually sits

Not in the code, and not in the compliance posture. In three places that no test suite can reach.

1. Nobody outside the build has used it

Nine synthetic households exercise the engine thoroughly and were chosen to include the awkward cases — a household in deficit, a freelancer with lumpy income, a 58-year-old badly overweight equity, someone who refuses equity entirely, a high earner who is nonetheless overspending. But nobody has yet abandoned the wizard at step four and told us why. Every completion-rate decision in Phase 2 is reasoned, and every one of them is a hypothesis.

2. The advice is honest in a way that may be unwelcome

Rohan's plan tells a 29-year-old that all three of his goals are underfunded and that he should surrender a policy he was sold. That is correct, and it is the point of the product. Whether people return to a tool that tells them the truth about their retirement is a real question, and the answer is not in the repository.

3. Data ages faster than anyone remembers to check

Two files carry dated external figures: the assumptions module (tax slabs, which change at each Budget; small-savings rates, which change quarterly) and the insurance premium tables, stamped with their source date. Premiums are subtracted from the surplus *before* anything is invested, so an error there distorts the entire plan rather than one line of it. An invariant catches gross drift; nothing catches gentle staleness. Both are due for their annual re-check.
